

Mathematical Methods for Artificial Intelligence Lab 4

Vytautas Kraujalis

2021-11-20

Contents

1	Gaussian Process	1
1.1	Reading Data	1
1.2	Required packages	2
1.3	Parallel processing	2
1.4	EDA, first look at the dataset	2
1.5	Fitting models	10
1.5.1	Gaussian Process - Variational Bayesian Multinomial Probit Regression	10
1.5.2	Gaussian Process with Radial Basis Function Kernel	12
1.5.3	Support Vector Machines with Linear Kernel	15
1.5.4	Support Vector Machines with Radial Basis Function Kernel	18
1.6	Comparison of models	20
2	SVM dataset	21
2.1	Required packages	21
2.2	Reading Data	22
2.3	EDA, first look at the dataset	22
2.4	Fitting models	26
2.4.1	SVM Novelty Detenction - Linear Kernel	26
2.4.2	SVM Novelty Detenction - Radial Basis Kernel	28
2.5	Conclusion	29

1 Gaussian Process

1.1 Reading Data

```
set.seed(123)

data_original <- read.csv("satellite_train.csv")
```

1.2 Required packages

```
library(SmartEDA)
library(dplyr)
library(ggplot2)
library(caret)

# Function for 6 class accuracy from confusion matrix
classAcc <- function(confusionMatrix) {
  class1 <- round(confusionMatrix$table[1, 1] / sum(confusionMatrix$table[, 1]) * 100, 1)
  class2 <- round(confusionMatrix$table[2, 2] / sum(confusionMatrix$table[, 2]) * 100, 1)
  class3 <- round(confusionMatrix$table[3, 3] / sum(confusionMatrix$table[, 3]) * 100, 1)
  class4 <- round(confusionMatrix$table[4, 4] / sum(confusionMatrix$table[, 4]) * 100, 1)
  class5 <- round(confusionMatrix$table[5, 5] / sum(confusionMatrix$table[, 5]) * 100, 1)
  class6 <- round(confusionMatrix$table[6, 6] / sum(confusionMatrix$table[, 6]) * 100, 1)
  acc <- c(class1, class2, class3, class4, class5, class6)
  names(acc) <- colnames(confusionMatrix$table)
  return(acc)
}
```

1.3 Parallel processing

```
library(parallel)
no_cores <- detectCores() - 1
library(doParallel)
cl <- makePSOCKcluster(no_cores)
registerDoParallel(cl)
```

1.4 EDA, first look at the dataset

```
ExpData(data_original, type=1)
```

##	Descriptions	Value
## 1	Sample size (nrow)	4435
## 2	No. of variables (ncol)	37
## 3	No. of numeric/integer variables	37
## 4	No. of factor variables	0
## 5	No. of text variables	0
## 6	No. of logical variables	0
## 7	No. of identifier variables	0
## 8	No. of date variables	0
## 9	No. of zero variance variables (uniform)	0
## 10	%. of variables having complete cases	100% (37)
## 11	%. of variables having >0% and <50% missing cases	0% (0)
## 12	%. of variables having >=50% and <90% missing cases	0% (0)
## 13	%. of variables having >=90% missing cases	0% (0)

We have a dataset of 4435 observations with 37 variables, all of the variables are numerical. All variables have no missing values. Let's change the response variable to factor type.

```
data <- data_original %>%
  mutate(V37 = as.factor(V37)) %>%
  rename(Target = V37)
```

```
data <- data %>%
  mutate(
    Target = as.factor(case_when(
      Target == "1" ~ "Red Soil",
      Target == "2" ~ "Cotton Crop",
      Target == "3" ~ "Grey Soil",
      Target == "4" ~ "Damp Grey Soil",
      Target == "5" ~ "Soil With Vegetation Stubble",
      Target == "7" ~ "Very Damp Grey Soil",
      TRUE ~ "ERROR"
    ))
  )
```

Let's look at the target variable frequencies

```
data %>%
  group_by(Target) %>%
  summarise(n = n()) %>%
  mutate(n_prop = round(n / sum(n) * 100, 2))
```

```
## # A tibble: 6 x 3
##   Target                n n_prop
##   <fct>              <int> <dbl>
## 1 Cotton Crop         479  10.8
## 2 Damp Grey Soil     415   9.36
## 3 Grey Soil          961  21.7
## 4 Red Soil          1072  24.2
## 5 Soil With Vegetation Stubble 470  10.6
## 6 Very Damp Grey Soil 1038  23.4
```

Our target variable has 6 classes, the smallest class has 415 (9.4%) observations while the biggest class has 1072 (24.2%) observations.

Let's look at descriptive statistics of each variable:

```
ExpNumStat(data, by = "A", round = 2, gp = "Target") %>%
  select(Vname, min, max, mean, median, SD)
```

```
##   Vname min max  mean median   SD
## 1    V1  40 104 69.47    68 13.65
## 10   V10  27 130 83.13    85 22.81
## 11   V11  50 145 98.97   101 16.68
## 12   V12  29 157 82.41    81 18.84
## 13   V13  40 102 69.37    68 13.66
## 14   V14  27 131 83.73    87 22.79
## 15   V15  53 145 99.41   101 16.69
## 16   V16  33 151 82.65    81 18.80
## 17   V17  40 104 69.13    68 13.56
```

```
## 18 V18 27 130 83.43      85 22.82
## 19 V19 56 139 99.24     101 16.73
## 2  V2 27 137 83.86      87 22.73
## 20 V20 34 157 82.62      81 18.84
## 21 V21 39 104 68.92      67 13.47
## 22 V22 27 130 83.14      85 22.83
## 23 V23 50 140 99.00     100 16.72
## 24 V24 29 154 82.48      81 18.92
## 25 V25 40 104 69.25      68 13.65
## 26 V26 27 131 83.67      85 22.77
## 27 V27 53 140 99.32     101 16.70
## 28 V28 34 154 82.67      81 18.88
## 29 V29 39 104 69.03      68 13.53
## 3  V3 56 140 99.32     101 16.67
## 30 V30 27 128 83.43      85 22.81
## 31 V31 50 145 99.18     101 16.74
## 32 V32 29 157 82.64      81 18.94
## 33 V33 40 104 68.80      67 13.44
## 34 V34 27 130 83.15      85 22.76
## 35 V35 50 145 99.06     100 16.66
## 36 V36 29 157 82.58      81 18.90
## 4  V4 33 154 82.56      83 18.70
## 5  V5 40 102 69.21      68 13.55
## 6  V6 27 137 83.50      85 22.81
## 7  V7 50 145 99.17     101 16.63
## 8  V8 29 157 82.48      81 18.71
## 9  V9 40 104 68.96      67 13.50
```

Nothing seems unordinary.

We should look at the correlation between variables

```
# Correlation

corr_simple <- function(df,sig=0.5){
  corr <- cor(df)
  #prepare to drop duplicates and correlations of 1
  corr[lower.tri(corr,diag=TRUE)] <- NA
  #drop perfect correlations
  corr[corr == 1] <- NA
  #turn into a 3-column table
  corr <- as.data.frame(as.table(corr))
  #remove the NA values from above
  corr <- na.omit(corr)
  #select significant values
  corr <- subset(corr, abs(Freq) > sig)
  #sort by highest correlation
  corr <- corr[order(-abs(corr$Freq)),]
  return(corr)
}

correlation_matrix = cor(data %>% select(-Target))

length(findCorrelation(correlation_matrix, cutoff = 0.99))
```

```
## [1] 0
```

```
length(findCorrelation(correlation_matrix, cutoff = 0.95))
```

```
## [1] 11
```

```
length(findCorrelation(correlation_matrix, cutoff = 0.9))
```


```
## [1] 23
```

We have 0 variables with higher than .99 correlation. We have 11 variables with higher than .95 correlation. We have 23 variables with higher than .9 correlation.

We'll look through random scatter plots and see how our target variable is separated across. We'll save some interesting combinations for later. P.S. we'll use jitter() function to overcome observation overlap.

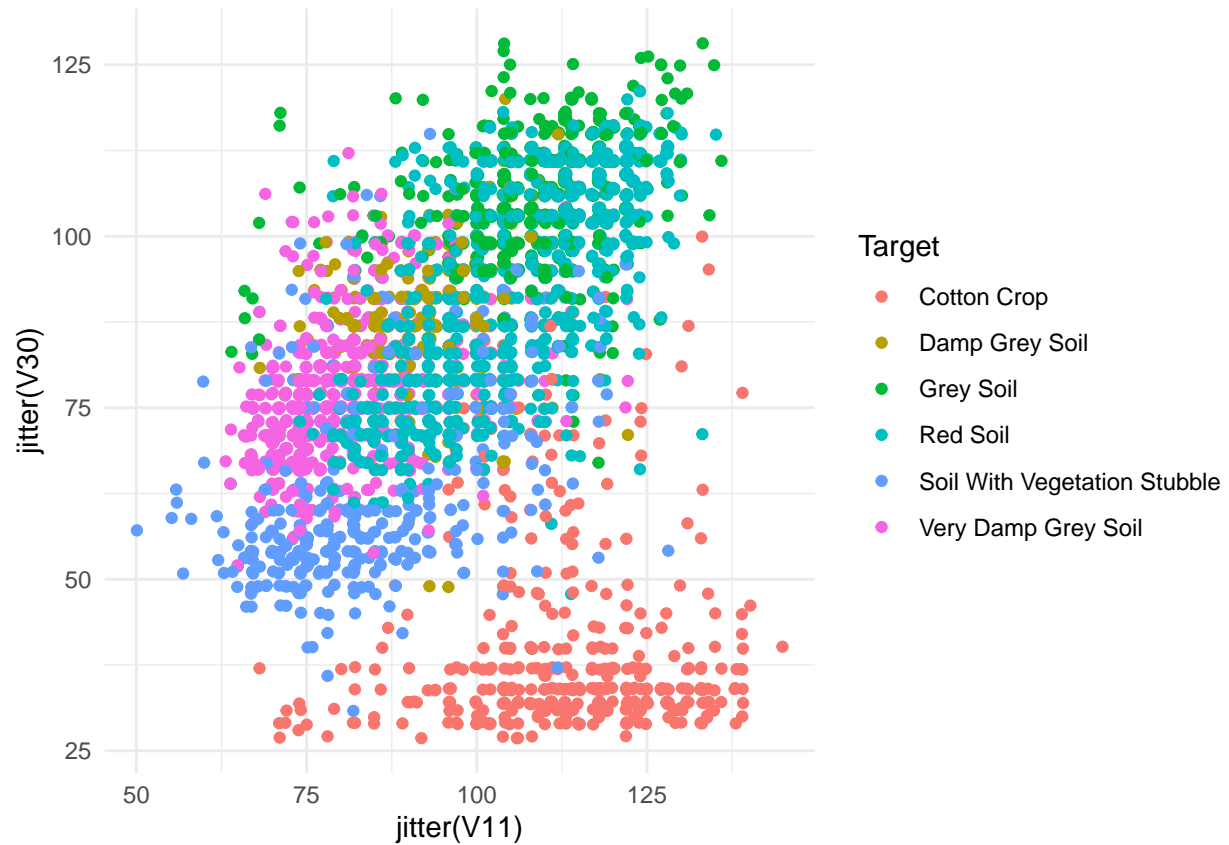
```
# data_plot <- data %>%  
#   select(sample(0:36, 1), sample(0:36, 1), Target)  
# data_plot_colnames <- colnames(data_plot)  
# colnames(data_plot) <- c("V_1", "V_2", "Target")  
# data_plot %>%  
#   ggplot(aes(x = jitter(V_1), y = jitter(V_2), color = Target)) +  
#   geom_point() +  
#   xlab(data_plot_colnames[1]) +  
#   ylab(data_plot_colnames[2])  
  
#combinations <- c("V2 - V3", "V30 - V11", "V7 - V3", "V3 - V24", "V31 - V32")
```

```
data %>%  
  ggplot(aes(x = jitter(V2), y = jitter(V3), color = Target)) +  
  geom_point() +  
  theme_minimal()
```




Combination of V2 and V3 features shows quite significant separation of “Cotton Crop” group, while “Damp Grey Soil” has almost no separation. All other classes could be visually separated.

```
data %>%
  ggplot(aes(x = jitter(V11), y = jitter(V30), color = Target)) +
  geom_point() +
  theme_minimal()
```




once again, the “Cotton Crop” group seems separatable quite good, but now there’s a mix of “Red Soil”, “Grey Soil” and “Damp Grey Soil” all in one cluster, which could be hard to distinguish.

```
data %>%
  ggplot(aes(x = jitter(V3), y = jitter(V7), color = Target)) +
  geom_point() +
  theme_minimal()
```



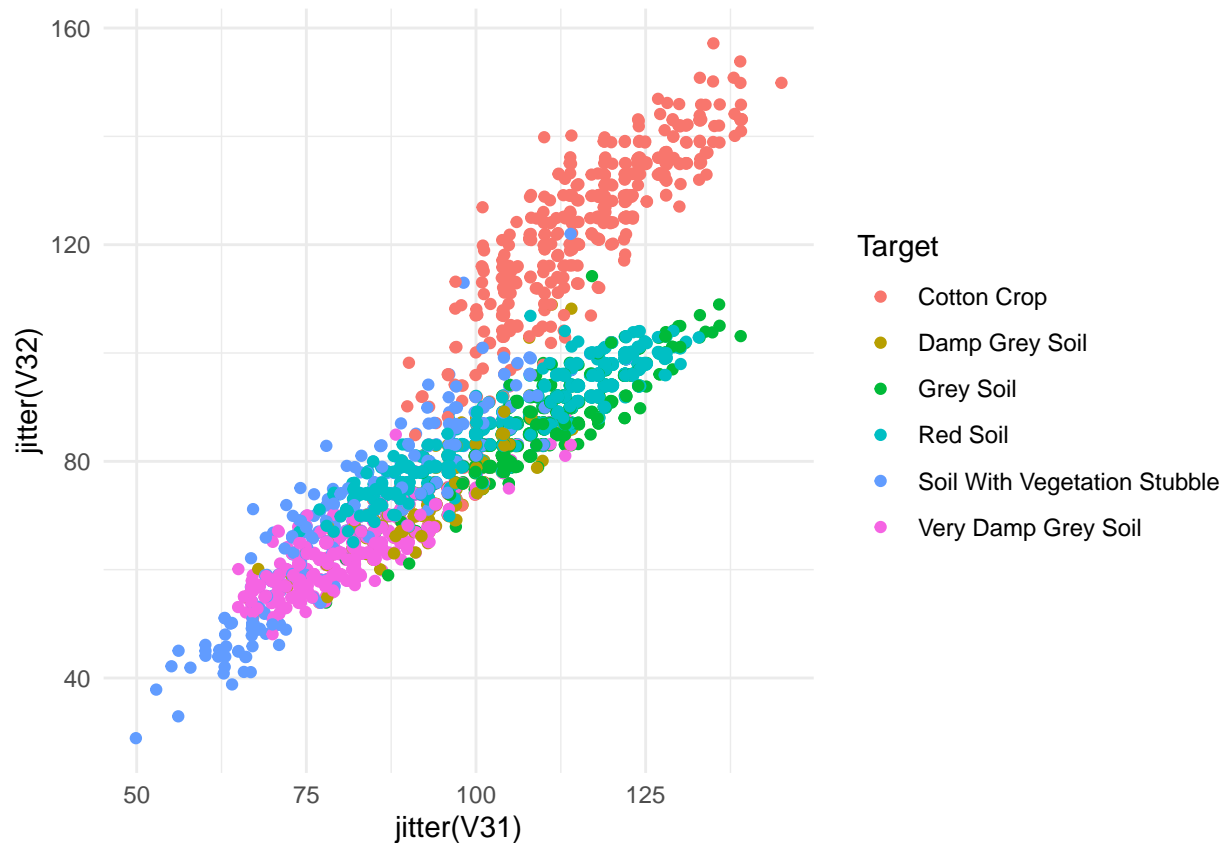
Looking at the combination of V3 and V7 features, we see a clear linear combination, but the target variable is mixed all over the place.

```
data %>%
  ggplot(aes(x = jitter(V3), y = jitter(V24), color = Target)) +
  geom_point() +
  theme_minimal()
```

Using a combination of V3 and V24 features we can separate “Cotton Crop” easily, but the rest of the classes won’t be separated so easily.

```
data %>%
  ggplot(aes(x = jitter(V31), y = jitter(V32), color = Target)) +
  geom_point() +
  theme_minimal()
```



Similar results as the plot before, but now the other classes could be separatable a bit better, only the “Damp Grey Soil” and “Soil With Vegetation Stubble” classes are not separatable that good.

It seems that our model could have a problem at detecting “Damp Grey Soil” class.

1.5 Fitting models

Tried fitting Linear Gaussian Process and Gaussian Process with Polynomial Kernel but both methods took too long to compute...

1.5.1 Gaussian Process - Variational Bayesian Multinomial Probit Regression

```
fitControl <- trainControl(
  method = "cv",
  number = 2,
  classProbs = TRUE,
  savePredictions="all",
  verboseIter = TRUE)

gp.vbmp.fit <- train(Target ~ ., data = data %>%
  mutate(Target = factor(Target, labels = make.names(levels(Target)))),
  method = "vbmpRadial",
  trControl = fitControl,
  preProcess=c("center", "scale","pca"))
```

```
## Aggregating results
## Fitting final model on full training set

resamp_vbmp = gp.vbmp.fit$pred[gp.vbmp.fit$pred$estimateTheta == gp.vbmp.fit$bestTune[1,1],]
confusion_matrix <- confusionMatrix(resamp_vbmp$pred, resamp_vbmp$obs)

confusion_matrix
```

```
## Confusion Matrix and Statistics
##
##               Reference
## Prediction   Cotton.Crop Damp.Grey.Soil Grey.Soil Red.Soil
## Cotton.Crop           461           5           2           2
## Damp.Grey.Soil         1          261          40           7
## Grey.Soil              3           72         889          16
## Red.Soil               0           5           7         1037
## Soil.With.Vegetation.Stubble 11          10           3          10
## Very.Damp.Grey.Soil     3           62          20           0
##
##               Reference
## Prediction   Soil.With.Vegetation.Stubble Very.Damp.Grey.Soil
## Cotton.Crop                                9           5
## Damp.Grey.Soil                             4          73
## Grey.Soil                                 4          31
## Red.Soil                                 21           0
## Soil.With.Vegetation.Stubble              394          30
## Very.Damp.Grey.Soil                       38         899
##
## Overall Statistics
##
##               Accuracy : 0.8886
##               95% CI : (0.879, 0.8977)
##               No Information Rate : 0.2417
##               P-Value [Acc > NIR] : < 2.2e-16
##
##               Kappa : 0.862
##
## Mcnemar's Test P-Value : NA
##
## Statistics by Class:
##
##               Class: Cotton.Crop Class: Damp.Grey.Soil Class: Grey.Soil
## Sensitivity           0.9624           0.62892           0.9251
## Specificity           0.9942           0.96891           0.9637
## Pos Pred Value        0.9525           0.67617           0.8759
## Neg Pred Value        0.9954           0.96197           0.9789
## Prevalence            0.1080           0.09357           0.2167
## Detection Rate        0.1039           0.05885           0.2005
## Detection Prevalence  0.1091           0.08703           0.2289
## Balanced Accuracy      0.9783           0.79891           0.9444
##
##               Class: Red.Soil Class: Soil.With.Vegetation.Stubble
## Sensitivity           0.9674           0.83830
## Specificity           0.9902           0.98386
## Pos Pred Value        0.9692           0.86026
## Neg Pred Value        0.9896           0.98089
```

```
## Prevalence          0.2417          0.10598
## Detection Rate      0.2338          0.08884
## Detection Prevalence 0.2413          0.10327
## Balanced Accuracy   0.9788          0.91108
##
## Class: Very.Damp.Grey.Soil
## Sensitivity         0.8661
## Specificity         0.9638
## Pos Pred Value      0.8796
## Neg Pred Value      0.9593
## Prevalence          0.2340
## Detection Rate      0.2027
## Detection Prevalence 0.2304
## Balanced Accuracy   0.9149
```

```
classAcc(confusion_matrix)
```

```
##          Cotton.Crop          Damp.Grey.Soil
##          96.2          62.9
##          Grey.Soil          Red.Soil
##          92.5          96.7
## Soil.With.Vegetation.Stubble Very.Damp.Grey.Soil
##          83.8          86.6
```

The overall accuracy of model is 89.7% which is much better than a random guess while NIR is 0.2417. The accuracies for each class shows what we predicted - the accuracy for “Damp Grey Soil” is only 60.2% while for other classes: + “Cotton Crop” - 95.6%, + “Grey Soil” - 94.1%, + “Red Soil” - 97.5%, + “Soil With Vegetation Stubble” - 84.3%, + Very Damp Grey Soil” - 89.1%.

1.5.2 Gaussian Process with Radial Basis Function Kernel

```
fitControl <- trainControl(
  method = "cv",
  number = 2,
  classProbs = TRUE,
  savePredictions="all",
  verboseIter = TRUE)


gpGrid = expand.grid(sigma = seq(0.01,0.2,0.01))

gp.fit <- train(Target ~ ., data = data %>%
  mutate(Target = factor(Target, labels = make.names(levels(Target)))),
  method = "gaussprRadial",
  trControl = fitControl,
  preProcess=c("center", "scale","pca"),
  tuneGrid = gpGrid)

## Aggregating results
## Selecting tuning parameters
## Fitting sigma = 0.18 on full training set
```

Let's look which sigma gives the best accuracy:

```
plot(gp.fit, metric = "Accuracy")
```



Best results are obtained when sigma is 0.15.

Using sigma = 0.15 let's look at the accuracies of each class and overall accuracy:

```
resamp_rb = gp.fit$pred[gp.fit$pred$sigma == gp.fit$bestTune[1,1],]
confusion_matrix <- confusionMatrix(resamp_rb$pred, resamp_rb$obs)
```

```
confusion_matrix
```

```
## Confusion Matrix and Statistics
```

```
##
```

```
##
```

```
## Prediction
```

```
## Cotton.Crop
```

```
## Damp.Grey.Soil
```

```
## Grey.Soil
```

```
## Red.Soil
```

```
## Soil.With.Vegetation.Stubble
```

```
## Very.Damp.Grey.Soil
```

```
##
```

```
## Prediction
```

```
## Cotton.Crop
```

```
Reference
```

```
Cotton.Crop Damp.Grey.Soil Grey.Soil Red.Soil
```

```
230 2 2 2
```

```
2 102 14 1
```

```
0 48 451 6
```

```
0 2 5 525
```

```
4 5 0 2
```

```
3 49 8 0
```

```
Reference
```

```
Soil.With.Vegetation.Stubble Very.Damp.Grey.Soil
```

```
2 1
```

```

## Damp.Grey.Soil 3 25
## Grey.Soil 1 11
## Red.Soil 15 0
## Soil.With.Vegetation.Stubble 173 5
## Very.Damp.Grey.Soil 41 477
##
## Overall Statistics
##
## Accuracy : 0.8832
## 95% CI : (0.8691, 0.8963)
## No Information Rate : 0.2418
## P-Value [Acc > NIR] : < 2.2e-16
##
## Kappa : 0.8543
##
## McNemar's Test P-Value : NA
##
## Statistics by Class:
##
## Class: Cotton.Crop Class: Damp.Grey.Soil Class: Grey.Soil
## Sensitivity 0.9623 0.49038 0.9396
## Specificity 0.9954 0.97760 0.9620
## Pos Pred Value 0.9623 0.69388 0.8723
## Neg Pred Value 0.9954 0.94879 0.9829
## Prevalence 0.1078 0.09382 0.2165
## Detection Rate 0.1037 0.04601 0.2034
## Detection Prevalence 0.1078 0.06631 0.2332
## Balanced Accuracy 0.9789 0.73399 0.9508
##
## Class: Red.Soil Class: Soil.With.Vegetation.Stubble
## Sensitivity 0.9795 0.73617
## Specificity 0.9869 0.99193
## Pos Pred Value 0.9598 0.91534
## Neg Pred Value 0.9934 0.96943
## Prevalence 0.2418 0.10600
## Detection Rate 0.2368 0.07803
## Detection Prevalence 0.2467 0.08525
## Balanced Accuracy 0.9832 0.86405
##
## Class: Very.Damp.Grey.Soil
## Sensitivity 0.9191
## Specificity 0.9405
## Pos Pred Value 0.8253
## Neg Pred Value 0.9744
## Prevalence 0.2341
## Detection Rate 0.2152
## Detection Prevalence 0.2607
## Balanced Accuracy 0.9298

```

```
classAcc(confusion_matrix)
```

```

## Cotton.Crop Damp.Grey.Soil
## 96.2 49.0
## Grey.Soil Red.Soil
## 94.0 97.9
## Soil.With.Vegetation.Stubble Very.Damp.Grey.Soil

```

##

73.6

91.9

The overall accuracy of model is 87.9% which is much better than a random guess while NIR is 0.2419. The accuracies for each class shows what we predicted - the accuracy for “Damp Grey Soil” is only 54.1% while for other classes: + “Cotton Crop” - 97.9%, + “Grey Soil” - 92.9%, + “Red Soil” - 97.0%, + “Soil With Vegetation Stubble” - 72.3%, + Very Damp Grey Soil” - 89.6%.

1.5.3 Support Vector Machines with Linear Kernel

```
fitControl <- trainControl(
  method = "cv",
  number = 10,
  classProbs = TRUE,
  savePredictions="all",
  verboseIter = TRUE)


svmGrid = expand.grid(C=seq(0.001,5,0.2))

svm.fit <- train(Target ~ ., data = data %>%
  mutate(Target = factor(Target, labels = make.names(levels(Target)))),
  method = "svmLinear",
  trControl = fitControl,
  preProcess=c("center", "scale","pca"),
  tuneGrid = svmGrid)
```

```
## Aggregating results
## Selecting tuning parameters
## Fitting C = 0.801 on full training set
```

Let's look which Cost gives us the best accuracy:

```
plot(svm.fit,metric = "Accuracy")
```



Best accuracy is obtained using Cost = 0.601.

Using Cost = 0.601 let's look at the accuracies:

```
resamp_svm = svm.fit$pred[svm.fit$pred$C==svm.fit$bestTune[1,1],]
confusion_matrix <- confusionMatrix(resamp_svm$pred, resamp_svm$obs)
confusion_matrix
```

Confusion Matrix and Statistics

##

##

Prediction

Cotton.Crop

Damp.Grey.Soil

Grey.Soil

Red.Soil

Soil.With.Vegetation.Stubble

Very.Damp.Grey.Soil

##

Prediction

Cotton.Crop

Damp.Grey.Soil

Grey.Soil

Red.Soil

Reference

Cotton.Crop Damp.Grey.Soil Grey.Soil Red.Soil

462 5 1 0

0 174 27 2

0 116 903 11

0 5 27 741

7 2 0 318

10 113 3 0

Reference

Soil.With.Vegetation.Stubble Very.Damp.Grey.Soil

43 1

9 121

2 25

18 0


```

## Soil.With.Vegetation.Stubble        60        1
## Very.Damp.Grey.Soil                 338       890
##
## Overall Statistics
##
##           Accuracy : 0.7283
##           95% CI : (0.7149, 0.7413)
##           No Information Rate : 0.2417
##           P-Value [Acc > NIR] : < 2.2e-16
##
##           Kappa : 0.6625
##
## Mcnemar's Test P-Value : NA
##
## Statistics by Class:
##
##           Class: Cotton.Crop Class: Damp.Grey.Soil Class: Grey.Soil
## Sensitivity                0.9645                0.41928            0.9396
## Specificity                0.9874                0.96045            0.9557
## Pos Pred Value             0.9023                0.52252            0.8543
## Neg Pred Value             0.9957                0.94125            0.9828
## Prevalence                 0.1080                0.09357            0.2167
## Detection Rate             0.1042                0.03923            0.2036
## Detection Prevalence       0.1154                0.07508            0.2383
## Balanced Accuracy           0.9759                0.68986            0.9477
##
##           Class: Red.Soil Class: Soil.With.Vegetation.Stubble
## Sensitivity                0.6912                0.12766
## Specificity                0.9851                0.91728
## Pos Pred Value             0.9368                0.15464
## Neg Pred Value             0.9092                0.89869
## Prevalence                 0.2417                0.10598
## Detection Rate             0.1671                0.01353
## Detection Prevalence       0.1784                0.08749
## Balanced Accuracy           0.8382                0.52247
##
##           Class: Very.Damp.Grey.Soil
## Sensitivity                0.8574
## Specificity                0.8634
## Pos Pred Value             0.6573
## Neg Pred Value             0.9520
## Prevalence                 0.2340
## Detection Rate             0.2007
## Detection Prevalence       0.3053
## Balanced Accuracy           0.8604

```

```
classAcc(confusion_matrix)
```

```

##           Cotton.Crop           Damp.Grey.Soil
##           96.5                41.9
##           Grey.Soil           Red.Soil
##           94.0                69.1
## Soil.With.Vegetation.Stubble Very.Damp.Grey.Soil
##           12.8                85.7

```


1.5.4 Support Vector Machines with Radial Basis Function Kernel

```
fitControl <- trainControl(  
  method = "cv",  
  number = 10,  
  classProbs = TRUE,  
  savePredictions="all",  
  verboseIter = TRUE)  
  
svmGrid = expand.grid(C=seq(0.001,5,0.2), sigma = seq(0.01, 0.3, 0.09))  
  
svm.radial.fit <- train(Target ~ ., data = data %>%  
  mutate(Target = factor(Target, labels = make.names(levels(Target)))),  
  method = "svmRadial",  
  trControl = fitControl,  
  preProcess=c("center", "scale","pca"),  
  tuneGrid = svmGrid)
```

```
## Aggregating results  
## Selecting tuning parameters  
## Fitting sigma = 0.1, C = 3.2 on full training set
```

Let's look at the accuracy using different parameters:

```
plot(svm.radial.fit, metric = "Accuracy")
```



Best accuracy is obtained using $\sigma = 0.1$, Cost = 2.8, let's use those parameters to look at the class accuracies:

```
resamp_svm_radial = svm.radial.fit$pred[svm.radial.fit$pred$C == svm.radial.fit$bestTune[1,2] & svm.radial.fit$bestTune[3,2]]
confusion_matrix <- confusionMatrix(resamp_svm_radial$pred, resamp_svm_radial$obs)
confusion_matrix
```

```
## Confusion Matrix and Statistics
##
##               Reference
## Prediction   Cotton.Crop Damp.Grey.Soil Grey.Soil Red.Soil
## Cotton.Crop           455           7           3           3
## Damp.Grey.Soil         6          199           5           1
## Grey.Soil              0          110          932           9
## Red.Soil               0           3          17          655
## Soil.With.Vegetation.Stubble 10           0           1          404
## Very.Damp.Grey.Soil      8          96           3           0
##
##               Reference
## Prediction   Soil.With.Vegetation.Stubble Very.Damp.Grey.Soil
## Cotton.Crop                        15           1
## Damp.Grey.Soil                      3          68
## Grey.Soil                          8          50
## Red.Soil                          19           0
## Soil.With.Vegetation.Stubble       150          1
## Very.Damp.Grey.Soil               275         918
##
## Overall Statistics
##
##               Accuracy : 0.7461
##               95% CI : (0.733, 0.7589)
##               No Information Rate : 0.2417
##               P-Value [Acc > NIR] : < 2.2e-16
##
##               Kappa : 0.6858
##
## McNemar's Test P-Value : NA
##
## Statistics by Class:
##
##               Class: Cotton.Crop Class: Damp.Grey.Soil Class: Grey.Soil
## Sensitivity           0.9499           0.47952           0.9698
## Specificity           0.9927           0.97935           0.9491
## Pos Pred Value        0.9401           0.70567           0.8404
## Neg Pred Value        0.9939           0.94799           0.9913
## Prevalence            0.1080           0.09357           0.2167
## Detection Rate        0.1026           0.04487           0.2101
## Detection Prevalence  0.1091           0.06359           0.2501
## Balanced Accuracy      0.9713           0.72944           0.9594
##
##               Class: Red.Soil Class: Soil.With.Vegetation.Stubble
## Sensitivity           0.6110           0.31915
## Specificity           0.9884           0.89508
## Pos Pred Value        0.9438           0.26502
```

```
## Neg Pred Value          0.8885          0.91729
## Prevalence              0.2417          0.10598
## Detection Rate          0.1477          0.03382
## Detection Prevalence    0.1565          0.12762
## Balanced Accuracy       0.7997          0.60712
##                          Class: Very.Damp.Grey.Soil
## Sensitivity              0.8844
## Specificity              0.8875
## Pos Pred Value          0.7062
## Neg Pred Value          0.9617
## Prevalence              0.2340
## Detection Rate          0.2070
## Detection Prevalence    0.2931
## Balanced Accuracy       0.8860
```

```
classAcc(confusion_matrix)
```

```
##          Cotton.Crop          Damp.Grey.Soil
##          95.0          48.0
##          Grey.Soil          Red.Soil
##          97.0          61.1
## Soil.With.Vegetation.Stubble    Very.Damp.Grey.Soil
##          31.9          88.4
```

Overall accuracy is 74.4% while the accuracies for each class are: + Cotton Crop - 95.8%, + Damp Grey Soil - 46.3%, + Grey Soil - 96.5%, + Red Soil - 60.5%, + Soil With Vegetation Stubble - 32.3%, + Very Damp Grey Soil - 88.6%

1.6 Comparison of models

```
accuracies <- data.frame() %>%
  bind_rows(
    classAcc(confusionMatrix(resamp_vbmp$pred, resamp_vbmp$obs)),
    classAcc(confusionMatrix(resamp_rb$pred, resamp_rb$obs)),
    classAcc(confusionMatrix(resamp_svm$pred, resamp_svm$obs)),
    classAcc(confusionMatrix(resamp_svm_radial$pred, resamp_svm_radial$obs))
  ) %>%
  bind_cols(
    model = c("GP - Variational Bayesian Multinomial Probit Reg.", "GP - Radial Basis", "SVM - Linear", "SVM - Radial Basis"),
    overall_accuracy = round(c(
      sum(diag(as.matrix(confusionMatrix(resamp_vbmp$pred, resamp_vbmp$obs)$table))) / sum(colSums(confusionMatrix(resamp_vbmp$pred, resamp_vbmp$obs)$table))),
      sum(diag(as.matrix(confusionMatrix(resamp_rb$pred, resamp_rb$obs)$table))) / sum(colSums(confusionMatrix(resamp_rb$pred, resamp_rb$obs)$table))),
      sum(diag(as.matrix(confusionMatrix(resamp_svm$pred, resamp_svm$obs)$table))) / sum(colSums(confusionMatrix(resamp_svm$pred, resamp_svm$obs)$table))),
      sum(diag(as.matrix(confusionMatrix(resamp_svm_radial$pred, resamp_svm_radial$obs)$table))) / sum(colSums(confusionMatrix(resamp_svm_radial$pred, resamp_svm_radial$obs)$table)))
    ), 2)
  ) %>%
  tibble::column_to_rownames(var = "model")

accuracies
```

```
##          Cotton.Crop Damp.Grey.Soil
```

## GP - Variational Bayesian Multinomial Probit Reg.	96.2	62.9
## GP - Radial Basis	96.2	49.0
## SVM - Linear Kernel	96.5	41.9
## SVM - Radial Kernel	95.0	48.0
##	Grey.Soil	Red.Soil
## GP - Variational Bayesian Multinomial Probit Reg.	92.5	96.7
## GP - Radial Basis	94.0	97.9
## SVM - Linear Kernel	94.0	69.1
## SVM - Radial Kernel	97.0	61.1
##	Soil.With.Vegetation.Stubble	
## GP - Variational Bayesian Multinomial Probit Reg.		83.8
## GP - Radial Basis		73.6
## SVM - Linear Kernel		12.8
## SVM - Radial Kernel		31.9
##	Very.Damp.Grey.Soil	
## GP - Variational Bayesian Multinomial Probit Reg.		86.6
## GP - Radial Basis		91.9
## SVM - Linear Kernel		85.7
## SVM - Radial Kernel		88.4
##	overall_accuracy	
## GP - Variational Bayesian Multinomial Probit Reg.		88.86
## GP - Radial Basis		88.32
## SVM - Linear Kernel		72.83
## SVM - Radial Kernel		74.61

Variational Bayesian Multinomial Probit Regression gave us the best overall accuracy of 89.7% while Gaussian Process with Radial Basis has 87.7% accuracy. Simple SVM with linear kernel gave us the worst overall accuracy of 72.5%.

Both Gaussian Process methods struggle with predicting “Damp Grey Soil” class, but that’s what we predicted from scatter plots. While SVM methods showed similar results predicting “Damp Grey Soil” class, but the SVM models had a very hard time predicting “Soil With Vegetation Stubble” while GP methods didnt had a very hard time predicting this class.

2 SVM dataset

2.1 Required packages

```
library(SmartEDA)
library(dplyr)
library(ggplot2)
library(caret)
library(R.matlab)
library(kernlab)

# Function for 6 class accuracy from confusion matrix
classAcc <- function(confusionMatrix) {
  class1 <- round(confusionMatrix$table[1, 1] / sum(confusionMatrix$table[, 1]) * 100, 1)
  class2 <- round(confusionMatrix$table[2, 2] / sum(confusionMatrix$table[, 2]) * 100, 1)
  acc <- c(class1, class2 )
  names(acc) <- colnames(confusionMatrix$table)
```

```

    return(acc)
}

```

2.2 Reading Data

```

set.seed(123)

annthyroid <- readMat("annthyroid.mat") %>% as.data.frame()

```

2.3 EDA, first look at the dataset

```
ExpData(annthyroid, type=1)
```

##		Descriptions	Value
## 1		Sample size (nrow)	7200
## 2		No. of variables (ncol)	7
## 3		No. of numeric/interger variables	7
## 4		No. of factor variables	0
## 5		No. of text variables	0
## 6		No. of logical variables	0
## 7		No. of identifier variables	0
## 8		No. of date variables	0
## 9		No. of zero variance variables (uniform)	0
## 10		%. of variables having complete cases	100% (7)
## 11	%. of variables having >0% and <50% missing cases		0% (0)
## 12	%. of variables having >=50% and <90% missing cases		0% (0)
## 13	%. of variables having >=90% missing cases		0% (0)

We have a dataset of 7200 observations, all 7 variables are of numeric. None of the columns have missing values.

Let's look at the target variable frequencies

```

annthyroid %>%
  group_by(y) %>%
  summarise(n = n()) %>%
  mutate(n_prop = round(n / sum(n) * 100, 2))

```

```

## # A tibble: 2 x 3
##       y     n n_prop
##   <dbl> <int> <dbl>
## 1     0  6666  92.6
## 2     1   534   7.42

```

It's a two-class problem, there's a huge class imbalance of 6666 (92.5%) / 534 (7.5%). In this lab, we are not going to try to address this problem.

Let's look at descriptive statistics of each variable:

```
ExpNumStat(annthyroid,by ="A",round= 2, gp = "y") %>%
  select(Vname, min, max, mean, median, SD)
```

```
##   Vname  min  max mean median   SD
## 1   X.1 0.01 0.97 0.52   0.55 0.19
## 2   X.2 0.00 0.53 0.00   0.00 0.02
## 3   X.3 0.00 0.18 0.02   0.02 0.01
## 4   X.4 0.00 0.60 0.11   0.11 0.04
## 5   X.5 0.02 0.23 0.10   0.10 0.02
## 6   X.6 0.00 0.64 0.11   0.11 0.04
```

Nothing seems unordinary.

We should look at the correlation between variables

```
correlation_matrix = cor(annthyroid %>% select(-y))

length(findCorrelation(correlation_matrix, cutoff = 0.99))
```

```
## [1] 0
```

```
length(findCorrelation(correlation_matrix, cutoff = 0.95))
```

```
## [1] 0
```

```
length(findCorrelation(correlation_matrix, cutoff = 0.9))
```

```
## [1] 0
```


We don't have any variables with higher correlation than 0.9.

We'll look through random scatter plots and see how our target variable is separated across. We'll save some interesting combinations for later. P.S. we'll use jitter() function to overcome observation overlap.

```
# data_plot <- annthyroid %>%
#   select(sample(0:6, 1), sample(0:6, 1), y) %>%
#   mutate(y = as.factor(y))
# data_plot_colnames <- colnames(data_plot)
# colnames(data_plot) <- c("V_1", "V_2", "y")
# data_plot %>%
#   ggplot(aes(x = jitter(V_1), y = jitter(V_2), color = y)) +
#   geom_point() +
#   xlab(data_plot_colnames[1]) +
#   ylab(data_plot_colnames[2])


#combinations <- c("X.3 - X.6", "X.2 - X.3", "X.5 - X.3")
```

```
annthyroid %>%
  ggplot(aes(x = jitter(X.3), y = jitter(X.6), color = as.factor(y))) +
  geom_point() +
  theme_minimal()
```




Combination of X.3 and X.6 features shows a very small separation between the 2 features. Seems like class 1 lies mostly in the range of X.6 [0 - 0.08] and X.3 [0 - 0.025].

```
annthyroid %>%  
  ggplot(aes(x = jitter(X.2), y = jitter(X.3), color = as.factor(y))) +  
  geom_point() +  
  theme_minimal()
```

This time, the class separation is much better, when the value of X.2 is greater than 0, there's a high chance the target class will be "1".

```
annthyroid %>%  
  ggplot(aes(x = jitter(X.5), y = jitter(X.3), color = as.factor(y))) +  
  geom_point() +  
  theme_minimal()
```



A combination of X.3 and X.5 almost shows us no class separation at all.

2.4 Fitting models

2.4.1 SVM Novelty Detention - Linear Kernel

```
fold_ids = createFolds(anthyroid$y, k = 5, list = TRUE, returnTrain = FALSE)

nu_list = seq(0.01,0.5,0.01)

collect_CV = array(0, dim = c(length(nu_list),5,5))

for(k in 1:5){
  for( i in 1:length(nu_list)){
    anomaly = ksvm(y ~ .,
                  anthyroid[-fold_ids[[k]], ],
                  kernel = "vanilladot",
                  type = 'one-svc',
                  nu = nu_list[i], kpar = list())
    y_true = anthyroid$y[fold_ids[[k]]]
    y_pred = 1 - 1*(predict(anomaly, anthyroid[fold_ids[[k]], ]))
    confMat = confusionMatrix(table(y_true = y_true, y_pred = y_pred))
    collect_CV[i,1,k] = confMat$overall[1] #overall accuracy
    collect_CV[i,2,k] = confMat$overall[2] #Kappa
  }
}
```

```

collect_CV[i,3,k] = confMat$overall[5] #NIR rate
collect_CV[i,4,k] = confMat$byClass[1] #Normal class (0) acc
collect_CV[i,5,k] = confMat$byClass[2] #Anomaly class (1) acc

}
print(paste("Done with fold:", k))
}

```

```

## [1] "Done with fold: 1"
## [1] "Done with fold: 2"
## [1] "Done with fold: 3"
## [1] "Done with fold: 4"
## [1] "Done with fold: 5"

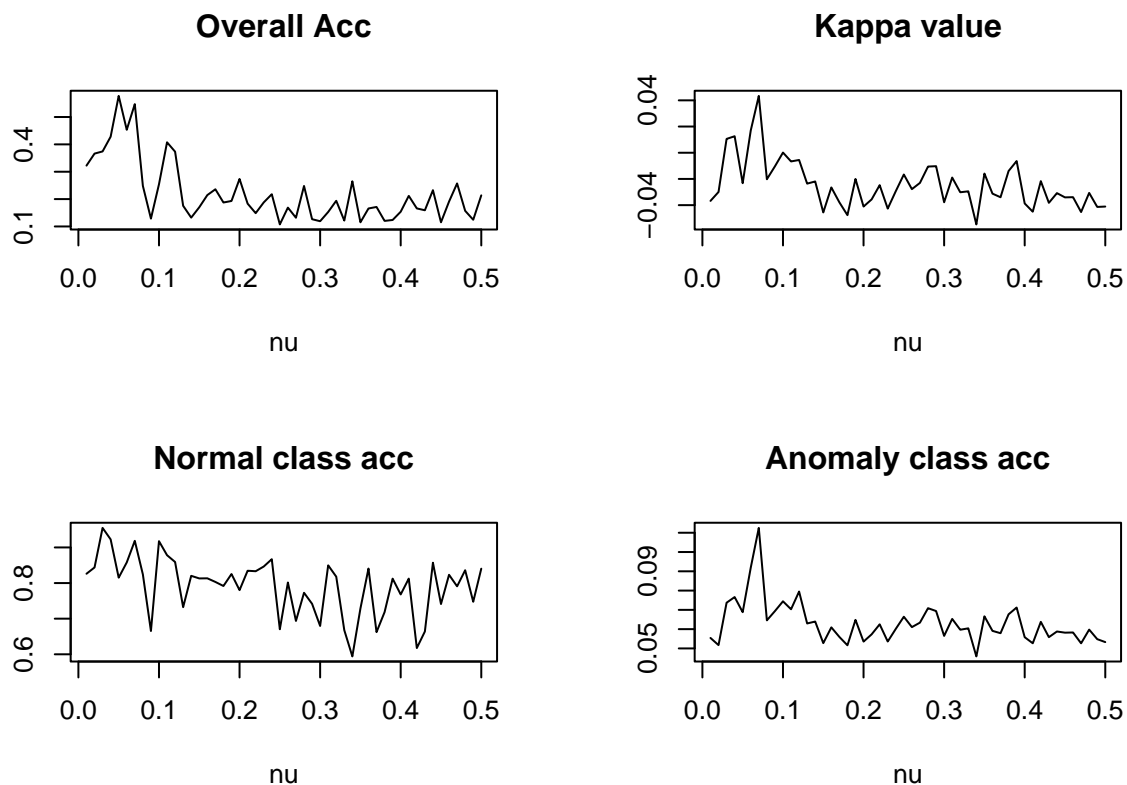
```

```
CV = apply(collect_CV, c(1,2), mean)
```

```

par(mfrow=c(2,2))
plot(nu_list,CV[,1],main="Overall Acc",xlab="nu",ylab="", type = "l")
plot(nu_list,CV[,2],main="Kappa value",xlab="nu",ylab="", type = "l")
plot(nu_list,CV[,4],main="Normal class acc",xlab="nu",ylab="", type = "l")
plot(nu_list,CV[,5],main="Anomaly class acc",xlab="nu",ylab="", type = "l")

```



Accuracy for the anomaly class is really small, no matter what ν value we are choosing.

2.4.2 SVM Novelty Detenction - Radial Basis Kernel

```
fold_ids = createFolds(annthyroid$y, k = 5, list = TRUE, returnTrain = FALSE)

nu_list = seq(0.01,0.5,0.01)

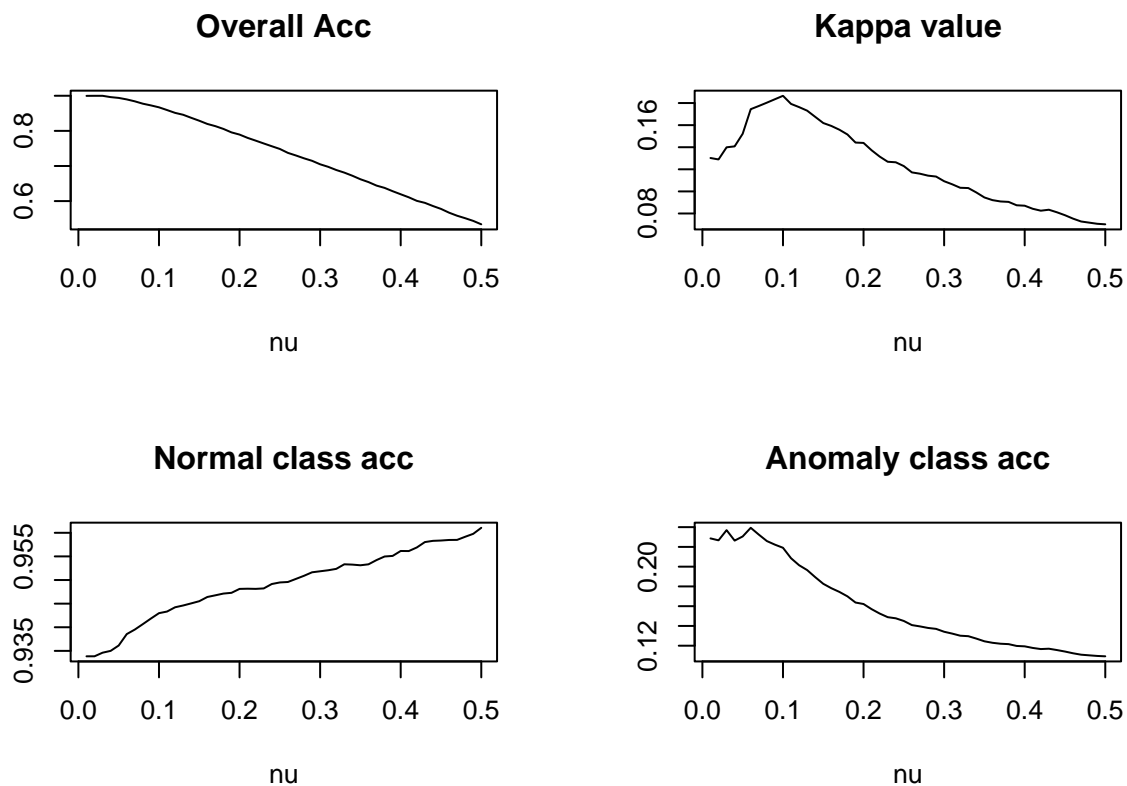
collect_CV = array(0, dim = c(length(nu_list),5,5))

for(k in 1:5){
  for( i in 1:length(nu_list)){
    anomaly = ksvm(y ~ .,
                  annthyroid[-fold_ids[[k]], ],
                  kernel = "rbfdot",
                  type = 'one-svc',
                  nu = nu_list[i], kpar = "automatic")
    y_true = annthyroid$y[fold_ids[[k]]]
    y_pred = 1 - 1*(predict(anomaly, annthyroid[fold_ids[[k]], ]))
    confMat = confusionMatrix(table(y_true = y_true, y_pred = y_pred))
    collect_CV[i,1,k] = confMat$overall[1] #overall accuracy
    collect_CV[i,2,k] = confMat$overall[2] #Kappa
    collect_CV[i,3,k] = confMat$overall[5] #NIR rate
    collect_CV[i,4,k] = confMat$byClass[1] #Normal class (0) acc
    collect_CV[i,5,k] = confMat$byClass[2] #Anomaly class (1) acc
  }
  print(paste("Done with fold:", k))
}
```

```
## [1] "Done with fold: 1"
## [1] "Done with fold: 2"
## [1] "Done with fold: 3"
## [1] "Done with fold: 4"
## [1] "Done with fold: 5"
```

```
CV = apply(collect_CV, c(1,2), mean)
```

```
par(mfrow=c(2,2))
plot(nu_list,CV[,1],main="Overall Acc",xlab="nu",ylab="", type = "l")
plot(nu_list,CV[,2],main="Kappa value",xlab="nu",ylab="", type = "l")
plot(nu_list,CV[,4],main="Normal class acc",xlab="nu",ylab="", type = "l")
plot(nu_list,CV[,5],main="Anomaly class acc",xlab="nu",ylab="", type = "l")
```



The higher the value of ν , the lower accuracy for the anomaly class. Highest accuracy of the anomaly class is reached using $\nu \sim 0.05$ and the accuracy is $\sim 24\%$

2.5 Conclusion

Seems like using Radial Basis Kernel gives much better prediction of anomaly class ($\sim 24\%$) while the linear kernel reached only 8% accuracy for the anomaly class